

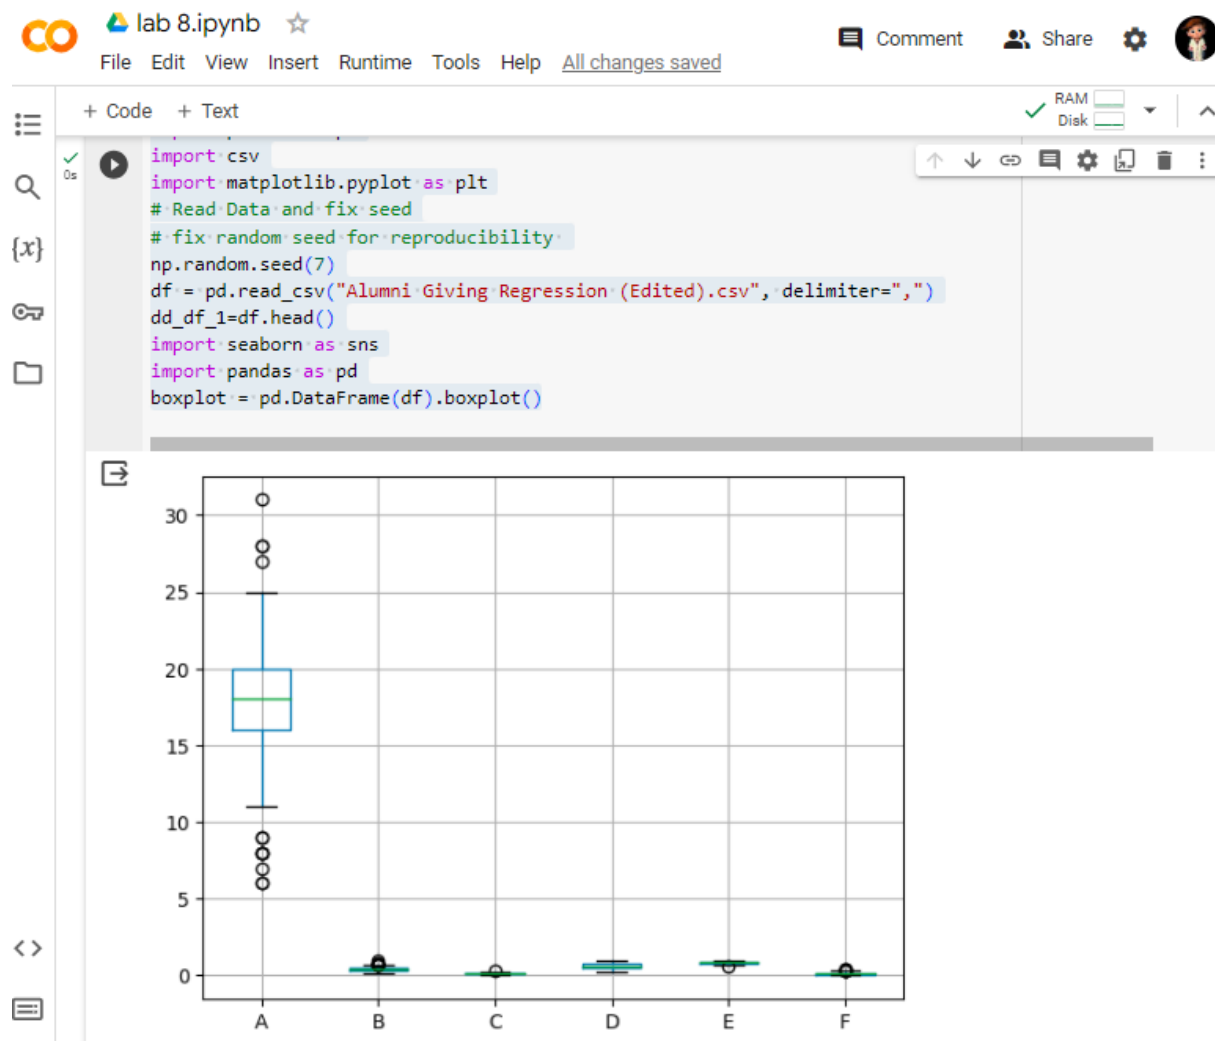
IN LAB

LAB TASK 1 :

Imports necessary libraries for machine learning and data visualization.

Reads a CSV file into a DataFrame and sets a random seed for reproducibility.

```
# Importing libraries needed
# Note that keras is generally used for deep learning as well from
keras.models import Sequential
from keras.layers import Dense, Dropout
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
import numpy as np
from sklearn import linear_model
from sklearn import preprocessing
from sklearn import tree
from sklearn.ensemble import RandomForestRegressor,
GradientBoostingRegressor
import pandas as pd
import csv
import matplotlib.pyplot as plt
# Read Data and fix seed
# fix random seed for reproducibility
np.random.seed(7)
df = pd.read_csv("Alumni Giving Regression (Edited).csv",
delimiter=",")
dd_df_1=df.head()
import seaborn as sns
import pandas as pd
boxplot = pd.DataFrame(df).boxplot()
```

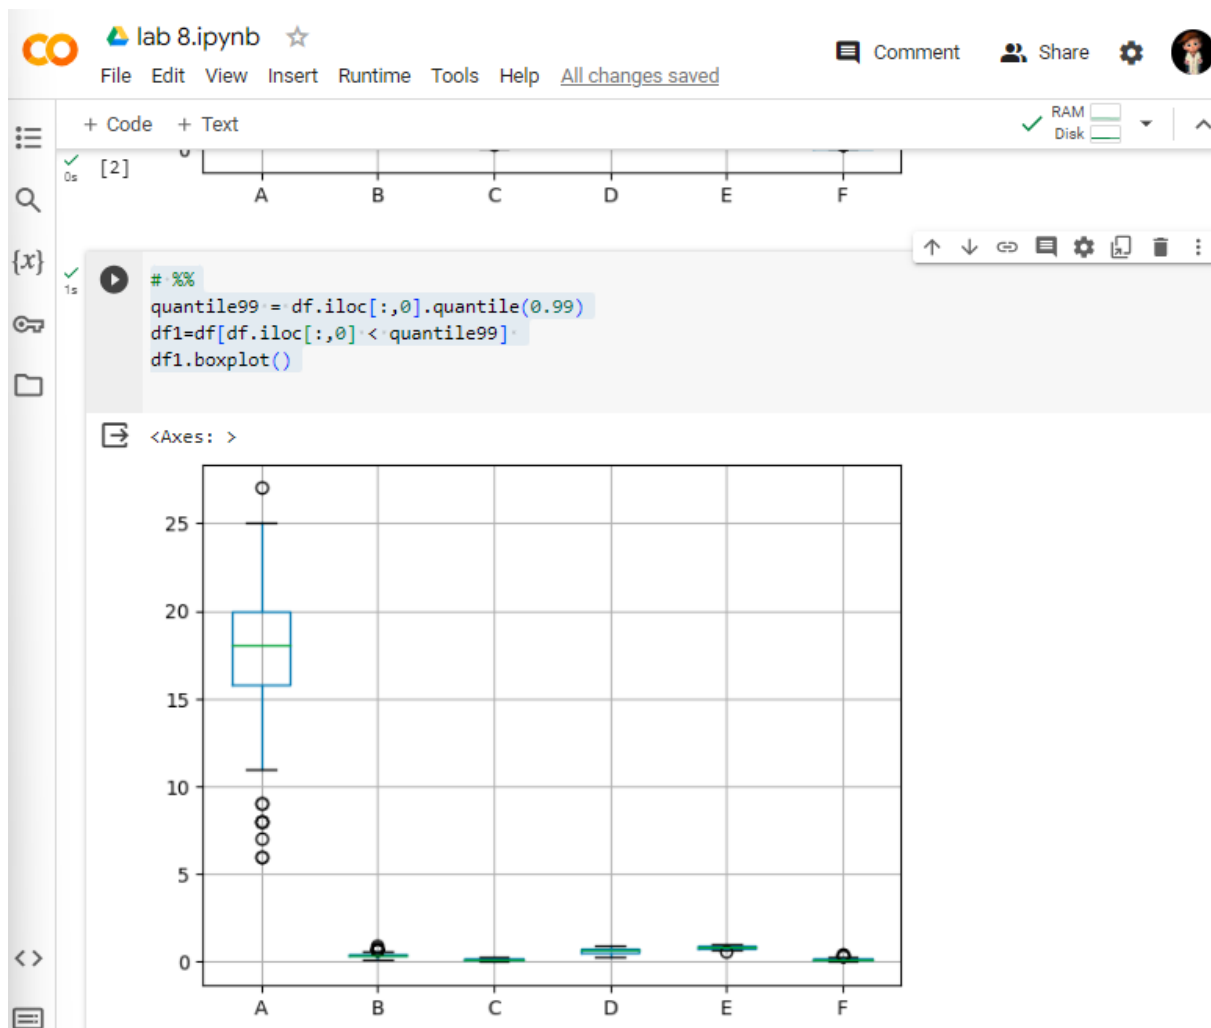


LAB TASK 2:

Filters data points above the 99th percentile, creating a new DataFrame (df1).

Visualizes the filtered data using a boxplot.

```
# %%
quantile99 = df.iloc[:,0].quantile(0.99)
df1=df[df.iloc[:,0] < quantile99]
df1.boxplot()
```

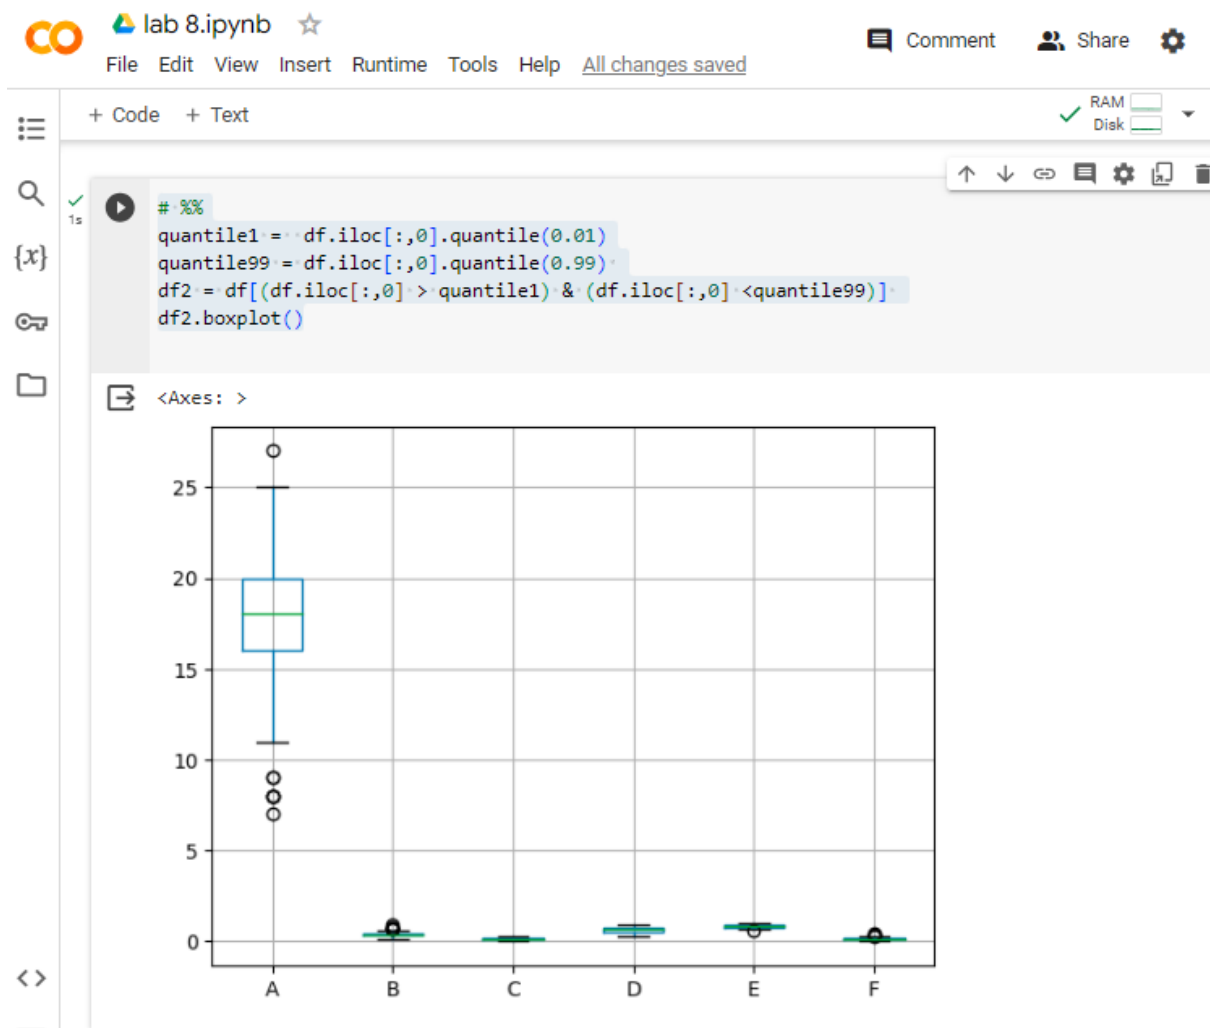


LAB TASK 3:

Filters data points between the 1st and 99th percentiles, creating a new DataFrame (df2).

Visualizes the filtered data using a boxplot.

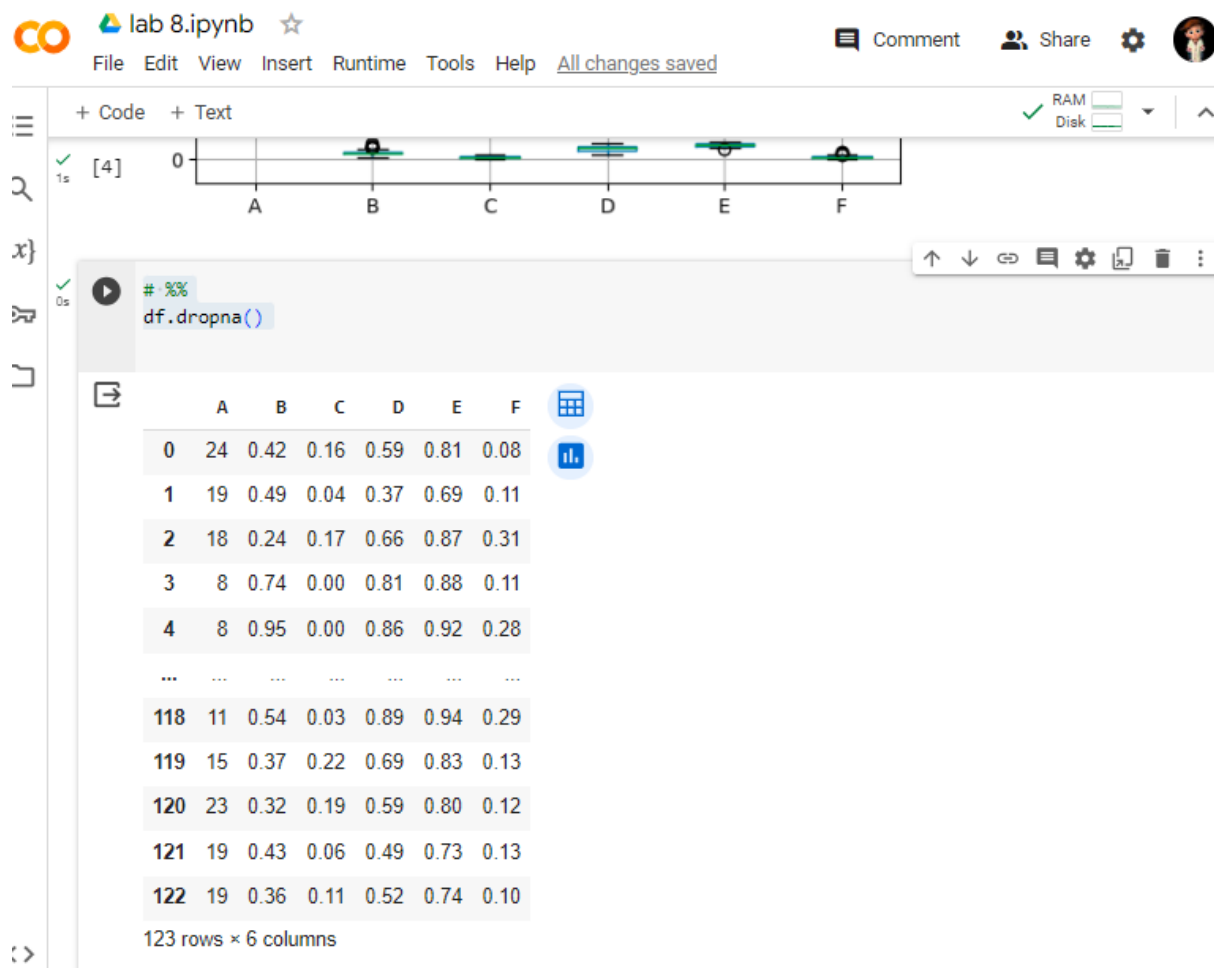
```
# %%
quantile1 = df.iloc[:,0].quantile(0.01)
quantile99 = df.iloc[:,0].quantile(0.99)
df2 = df[(df.iloc[:,0] > quantile1) & (df.iloc[:,0] < quantile99)]
df2.boxplot()
```



LAB TASK 4:

Drops rows with missing values from the DataFrame.

```
# %%  
df.dropna()
```



LAB TASK 5:

Drops missing values and separates features (X) and target variable (y).

Ranks features based on importance using a Random Forest Regressor.

```
# Assuming you have already loaded the data into the 'df' DataFrame

# Dropping missing values
df = df.dropna()

# Feature Ranking
X = df.iloc[:, 1:]
y = df.iloc[:, 0]

model3 = RandomForestRegressor()
model3.fit(X, y)

importances = model3.feature_importances_
std = np.std([tree.feature_importances_ for tree in
model3.estimators_], axis=0)
indices = np.argsort(importances)[-1:-1]
```

```
# Print the feature ranking
print("Feature ranking: ")
for f in range(X.shape[1]):
    print("%d. feature (Column index) %s (%f)" % (f + 1, indices[f],
importances[indices[f]]))
```

The screenshot shows a Jupyter Notebook titled 'lab 8.ipynb'. The code cell contains the following Python code:

```
model3 = RandomForestRegressor()
model3.fit(X, y)

importances = model3.feature_importances_
std = np.std([tree.feature_importances_ for tree in model3.estimators_], axis=0)
indices = np.argsort(importances)[::-1]

# Print the feature ranking
print("Feature ranking:")
for f in range(X.shape[1]):
    print("%d. feature (Column index) %s (%f)" % (f + 1, indices[f], importances[indices[f]]))
```

The output of the code is displayed below the code cell:

```
Feature ranking:
1. feature (Column index) 0 (0.547426)
2. feature (Column index) 2 (0.163201)
3. feature (Column index) 1 (0.121114)
4. feature (Column index) 3 (0.099430)
5. feature (Column index) 4 (0.068829)
```

LAB TASK 6

Selects top 3 important features based on the ranking.

Builds a linear regression model, evaluates on training and testing sets using RMSE.

```
# %%
indices_top3= indices[:3]
print(indices_top3)
dataset=df
df = pd.DataFrame(df)
Y_position = 5
TOP_N_FEATURE = 3
X = dataset.iloc[:, indices_top3]
Y = dataset.iloc[:,Y_position]
# create model
X_train, X_test, y_train, y_test = train_test_split(X, Y,
test_size=0.20, random_state=2020)
#Model 1 linear regression
model1 = linear_model.LinearRegression()
```

```

modell1.fit(X_train, y_train)
y_pred_train1 = modell1.predict(X_train)
#print("Regression")
#print("=====")

RMSE_train1 = mean_squared_error(y_train, y_pred_train1)
print("Regression TrainSet: RMSE {}".format(RMSE_train1))
#print("=====")
y_pred1 = modell1.predict(X_test)
RMSE_test1 = mean_squared_error(y_test, y_pred1)
print("Regression Testset: RMSE {}".format(RMSE_test1))
#print("=====")

```

The screenshot shows a Jupyter Notebook interface with a code cell containing the following Python code:

```

print(indices_top3)
dataset=df
df = pd.DataFrame(df)
Y_position = 5
TOP_N_FEATURE = 3
X = dataset.iloc[:, indices_top3]
Y = dataset.iloc[:, Y_position]
# create model
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.20, random_state=2020)
#Model 1 linear regression
modell1 = linear_model.LinearRegression()
modell1.fit(X_train, y_train)
y_pred_train1 = modell1.predict(X_train)
#print("Regression")
#print("=====")

RMSE_train1 = mean_squared_error(y_train, y_pred_train1)
print("Regression TrainSet: RMSE {}".format(RMSE_train1))
#print("=====")
y_pred1 = modell1.predict(X_test)
RMSE_test1 = mean_squared_error(y_test, y_pred1)
print("Regression Testset: RMSE {}".format(RMSE_test1))
#print("=====")

```

The output of the code is displayed in the console:

```

[0 2 1]
Regression TrainSet: RMSE 0.0036988847883733275
Regression Testset: RMSE 0.005388812554401423

```